

1、什么是GTID?

- 1、全局唯一，一个事务对应一个GTID
- 2、替代传统的binlog+pos复制；使用master_auto_position=1自动匹配GTID断点进行复制
- 3、MySQL5.6开始支持
- 4、在传统的主从复制中，slave端不用开启binlog；但是在GTID主从复制中，必须开启binlog
- 5、slave端在接受master的binlog时，会校验GTID值
- 6、为了保证主从数据的一致性，多线程同时执行一个GTID

2、组成

Master_UUID:序列号举例：ceb0ca3d-8366-11e8-ad2b-000c298b7c9a:1-5ceb0ca3d-8366-11e8-ad2b-000c298b7c9a其实就是master的uuid值；1-5是序列号，每次一个事务完成都会自增1，也就是说下一次为1-6。

3、工作原理

- 1、master更新数据时，会在事务前产生GTID，一同记录到binlog日志中。
- 2、slave端的i/o 线程将变更的binlog，写入到本地的relay log中。
- 3、sql线程从relay log中获取GTID，然后对比slave端的binlog是否有记录。
- 4、如果有记录，说明该GTID的事务已经执行，slave会忽略。
- 5、如果没有记录，slave就会从relay log中执行该GTID的事务，并记录到binlog。
- 6、在解析过程中会判断是否有主键，如果没有就用二级索引，如果没有就用全部扫描

master添加配置

```
[mysqld]
basedir = /usr/local/mysql
datadir = /usr/local/mysql/data
port = 12308
socket = /tmp/mysql.sock
character-set-server = utf8
log-error=/usr/local/mysql/mysql.log
pid-file=/usr/local/mysql/mysql.pid
lower_case_table_names=1
event_scheduler = 1

server-id=1
log-bin=/usr/local/mysql/binlog/mysql-bin.log
sql_mode =
"STRICT_TRANS_TABLES,NO_ZERO_IN_DATE,NO_ZERO_DATE,ERROR_FOR_DIVISION_BY_ZERO,NO_AUTO_CREATE_USER,NO_ENGINE_SUBSTITUTION"
log_bin_trust_function_creators=1

gtid_mode = on #开启GTID模式
enforce_gtid_consistency = on #使用GTID模式复制时，需要开启参数，用来保证数据的一致性
强制gtid复制
binlog_format = row #binlog格式为row
log-slave-updates = 1 #决定SLAVE从Master接收到更新且执行是否记录到SLAVE的binlog中
skip_slave_start = 1 #当SLAVE数据库启动的时候，SLAVE不会启动复制
```

```
auto-increment-increment = 2
auto-increment-offset = 1
```

slave添加配置

```
[mysqld]
basedir = /usr/local/mysql
datadir = /usr/local/mysql/data
port = 12308
socket = /tmp/mysql.sock
character-set-server = utf8
log-error=/usr/local/mysql/mysqlld.log
pid-file=/usr/local/mysql/mysqlld.pid
lower_case_table_names=1
event_scheduler = 1

slow_query_log=ON
slow_query_log_file=/usr/local/mysql/mysql-slow.log
long_query_time=1

server-id=2
log-bin=/usr/local/mysql/binlog/mysql-bin.log
sql_mode =
"STRICT_TRANS_TABLES,NO_ZERO_IN_DATE,NO_ZERO_DATE,ERROR_FOR_DIVISION_BY_ZERO,NO_AUTO_CREATE_USER,NO_ENGINE_SUBSTITUTION"
log_bin_trust_function_creators=1

gtid_mode = on #开启GTID模式
enforce_gtid_consistency = on #使用GTID模式复制时，需要开启参数，用来保证数据的一致性强制gtid复制
binlog_format = row #binlog格式为row
log-slave-updates = 1 #决定SLAVE从Master接收到更新且执行是否记录到SLAVE的binlog中
skip_slave_start = 1 #当SLAVE数据库启动的时候，SLAVE不会启动复制
auto-increment-increment = 2
auto-increment-offset = 1
```

server-id 值主从要不一致

master授权配置

```
mysql -uroot -p
mysql> grant replication client,replication slave on *.* to 'rep'@'172.60.51.%' identified by 'qwe123';
mysql> flush privileges;
```

slave配置同步

```
mysql -uroot -p
```

```
mysql> change master to master_host='172.60.51.91',  
master_user='rep',master_password='qwe123',master_port=12308,master_auto_position=1;  
mysql> start slave;
```

查看slave的状态

```
mysql> show slave status\G
```

```
***** 1. row *****  
Slave_IO_State: Waiting for master to send event  
Master_Host: 172.60.17.10  
Master_User: rep  
Master_Port: 12308  
Connect_Retry: 60  
Master_Log_File: mysql-bin.000005  
Read_Master_Log_Pos: 116849436  
Relay_Log_File: t02-relay-bin.000003  
Relay_Log_Pos: 116849649  
Relay_Master_Log_File: mysql-bin.000005  
Slave_IO_Running: Yes  
Slave_SQL_Running: Yes  
Replicate_Do_DB:  
Replicate_Ignore_DB:  
Replicate_Do_Table:  
Replicate_Ignore_Table:  
Replicate_Wild_Do_Table:  
Replicate_Wild_Ignore_Table:  
Last_Errno: 0  
Last_Error:  
Skip_Counter: 0  
Exec_Master_Log_Pos: 116849436  
Relay_Log_Space: 116849854  
Until_Condition: None  
Until_Log_File:  
Until_Log_Pos: 0  
Master_SSL_Allowed: No  
Master_SSL_CA_File:  
Master_SSL_CA_Path:  
Master_SSL_Cert:  
Master_SSL_Cipher:  
Master_SSL_Key:  
Seconds_Behind_Master: 0  
Master_SSL_Verify_Server_Cert: No  
Last_IO_Errno: 0  
Last_IO_Error:  
Last_SQL_Errno: 0  
Last_SQL_Error:  
Replicate_Ignore_Server_Ids:  
Master_Server_Id: 1  
Master_UUID: cd5d521b-b1c7-11e9-a1dc-000c2980d8b3  
Master_Info_File: /usr/local/mysql/data/master.info  
SQL_Delay: 0  
SQL_Remaining_Delay: NULL  
Slave_SQL_Running_State: Slave has read all relay log; waiting for more updates  
Master_Retry_Count: 86400  
Master_Bind:  
Last_IO_Error_Timestamp:  
Last_SQL_Error_Timestamp:  
Master_SSL_Crl:  
Master_SSL_Crlpath:
```

```
Retrieved_Gtid_Set: cd5d521b-b1c7-11e9-a1dc-000c2980d8b3:1-355
Executed_Gtid_Set: cd5d521b-b1c7-11e9-a1dc-000c2980d8b3:1-355
  Auto_Position: 1
Replicate_Rewrite_DB:
  Channel_Name:
  Master_TLS_Version:
1 row in set (0.00 sec)
```

```
Master_Host: 172.60.17.10
Master_User: rep
Master_Port: 12308
Connect_Retry: 60
Master_Log_File: mysql-bin.000005
Read_Master_Log_Pos: 116849436
Relay_Log_File: t02-relay-bin.000003
Relay_Log_Pos: 116849649
Relay_Master_Log_File: mysql-bin.000005
Slave_IO_Running: Yes
Slave_SQL_Running: Yes
Replicate_Do_DB:
```

```
Master Server Id: 1
Master_UUID: cd5d521b-b1c7-11e9-a1dc-000c2980d8b3
Master_Info_File: /usr/local/mysql/data/master.info
  SQL_Delay: 0
  SQL_Remaining_Delay: NULL
Slave_SQL_Running_State: Slave has read all relay log; waiting for more updates
Master_Retry_Count: 86400
Master_Bind:
Last_IO_Error_Timestamp:
Last_SQL_Error_Timestamp:
  Master_SSL_Crl:
  Master_SSL_Crlpath:
Retrieved_Gtid_Set: cd5d521b-b1c7-11e9-a1dc-000c2980d8b3:1-355
Executed_Gtid_Set: cd5d521b-b1c7-11e9-a1dc-000c2980d8b3:1-355
  Auto_Position: 1
Replicate_Rewrite_DB:
  Channel_Name:
  Master_TLS_Version:
1 row in set (0.00 sec)
```

Retrieved_Gtid_Set 表示slave从master接受的gtid set, 使用 reset slave 命令可以清空此项;

Executed_Gtid_Set 表示slave已执行的gtid set, 使用 reset master 命令可以清空此项。

Retrieved_Gtid_Set 和 Executed_Gtid_Set 必须为master 上 gtid set 的子集, 否则会报以下错误:

查看master状态

```
mysql> show master status;
+-----+-----+-----+-----+-----+
| File           | Position | Binlog_Do_DB | Binlog_Ignore_DB | Executed_Gtid_Set |
+-----+-----+-----+-----+-----+
| mysql-bin.000005 | 116849436 |              |                  | cd5d521b-b1c7-11e9-a1dc-000c2980d8b3:1-355 |
+-----+-----+-----+-----+-----+
1 row in set (0.00 sec)
```

对比slave端, Executed_Gtid_Set的值应该是一样的。

参考URL: <https://blog.5lcto.com/13434336/2178937>

master uuid: cd5d521b-b1c7-11e9-a1dc-000c2980d8b3

slave uuid: 08109ba1-b994-11e9-b0ad-000c290b47ec

其他:

NO. 1 从库有数据写入（即从库插入数据）

mysql> create database t2;

```
Until_Log_File:
Until_Log_Pos: 0
Master_SSL_Allowed: No
Master_SSL_CA_File:
Master_SSL_CA_Path:
Master_SSL_Cert:
Master_SSL_Cipher:
Master_SSL_Key:
Seconds_Behind_Master: 0
Master_SSL_Verify_Server_Cert: No
Last_IO_Errno: 0
Last_IO_Error:
Last_SQL_Errno: 0
Last_SQL_Error:
Replicate_Ignore_Server_Ids:
Master_Server_Id: 1
Master_UUID: cd5d521b-b1c7-11e9-a1dc-000c2980d8b3
Master_Info_File: /usr/local/mysql/data/master.info
SQL_Delay: 0
SQL_Remaining_Delay: NULL
Slave_SQL_Running_State: Slave has read all relay log; waiting for more updates
Master_Retry_Count: 86400
Master_Bind:
Last_IO_Error_Timestamp:
Last_SQL_Error_Timestamp:
Master_SSL_Crl:
Master_SSL_Crlpath:
Retrieved_Gtid_Set: cd5d521b-b1c7-11e9-a1dc-000c2980d8b3:1-356
Executed_Gtid_Set: 08109ba1-b994-11e9-b0ad-000c290b47ec:1-3,
cd5d521b-b1c7-11e9-a1dc-000c2980d8b3:1-356
Auto_Position: 1
Replicate_Rewrite_DB:
Channel_Name:
Master_TLS_Version:
1 row in set (0.00 sec)
```

可以看到已经执行的事务有来自主库的cd5d521b-b1c7-11e9-a1dc-000c2980d8b3:1-356,

也有从库自己写入的数据: 08109ba1-b994-11e9-b0ad-000c290b47ec:1-3。我们可以解析binlog看看

```
/usr/local/mysql/bin/mysqlbinlog -vv mysql-bin.000001 --include-gtids='08109ba1-b994-11e9-b0ad-000c290b47ec:1-3'
```

```

# at 116847597
#190809 10:38:34 server id 2  end_log_pos 116847683 CRC32 0x1a859e0d  Query  thread_id=17  exec_time=0  error_code=0
SET TIMESTAMP=1565318314/*!*/;
drop database t2
/*!*/;
# at 116847683
#190809 10:49:02 server id 2  end_log_pos 116847748 CRC32 0xe4d16d30  GTID    last_committed=358      sequence_number=359      rbr_only=no
SET @@SESSION.GTID_NEXT= '08109ba1-b994-11e9-b0ad-000c290b47ec:3'/*!*/;
# at 116847748
#190809 10:49:02 server id 2  end_log_pos 116847841 CRC32 0x57c445bd  Query  thread_id=17  exec_time=0  error_code=0
SET TIMESTAMP=1565318942/*!*/;
create database t2
/*!*/;
SET @@SESSION.GTID_NEXT= 'AUTOMATIC' /* added by mysqlbinlog */ /*!*/;
DELIMITER ;
# End of log file
/*!50003 SET COMPLETION_TYPE=@OLD_COMPLETION_TYPE*/;
/*!50530 SET @@SESSION.PSEUDO_SLAVE_MODE=0*/;

```

从binlog中可以清楚看到是从库进行了写入。